

1 Fake news: an algorithmic
2 perspective on fact-checking

3 Martijn B.J. Schouten
4 11295562

5 Bachelor thesis
6 Credits: 12 EC

7 Bachelor's degree Information Science

8 University of Amsterdam
9 Faculty of Science
10 Science Park 904
11 1098 XH Amsterdam

12 *Supervisor*
13 Dr. M. J. Marx

14 ILPS, IvI
15 Faculty of Science
16 University of Amsterdam
17 Science Park 904
18 1098 XH Amsterdam

19 2019-06

20

Abstract

21

22	Contents	
23	1 Introduction	4
24	2 Related Work	4
25	2.1 RQ1	4
26	2.2 RQ2	4
27	2.3 RQ3	5
28	3 Methodology	5
29	3.1 Description of the data	5
30	3.2 Data charts	5
31	3.3 Methods	5
32	3.3.1 RQ1	5
33	3.3.2 RQ2	5
34	4 Evaluation	5
35	5 Conclusions	5
36	5.1 Acknowledgements	5

1 Introduction

- Bevat je onderzoeksvraag (of vragen)
- Plaatst je vraag in de bestaande literatuur.

Je onderzoeksvraag is leidend voor je hele scriptie. Alles wat je doet moet uiteindelijk terug te voeren zijn op 1 doel: het beantwoorden van die vraag.

Typisch zal je het dan ook zo doen:

Mijn onderzoeksvraag is onderverdeeld in de volgende deelvragen:

RQ1 ... We beantwoorden deze vraag door het volgende te doen/ antwoord op de volgende vragen te vinden/ ...

1. Vragen op dit niveau kan je echt beantwoorden, en dat doe je in je Evaluatie sectie.

RQ2 ...

RQ3 ...

Je Evaluatie sectie bevat evenveel subsecties als je deelvragen hebt. En in elke sectie beantwoord je dan die deelvraag met behulp van de vragen op het onderste niveau.

In je conclusies kan je dan je hoofdvraag gaan beantwoorden op basis van al het eerder vergaarde bewijs.

Overview of thesis Hier geef je even kort weer wat in elke sectie staat.

2 Related Work

2.1 RQ1

Fake news as a term only caught public attention starting from the end of 2016, during the Presidential Elections of the United States [7].

2.2 RQ2

In the last couple of years, using transfer learning for natural language processing has given promisable results. The following sentence embeddings will be used to detect fake news:

- Bag of Words as a baseline for performance of non-pretrained embeddings;
- Facebook's InferSent [1];
- ELMo from the Allen Institute for Artificial Intelligence [5];
- OpenAI's GPT-2 [6];
- Transformer-XL [2];
- Microsoft's MT-DNN [4];
- and Google's BERT [3].

71 2.3 RQ3

72 Aligned with the original research on this dataset by Wang [8], the following
73 machine learning algorithms will be used to test the applicability of the abovementioned
74 embedding techniques:

- 75 • SVMs;
- 76 • Logistic regression;
- 77 • Bi-LSTMs;
- 78 • CNNs.

79 3 Methodology

80 3.1 Description of the data

81 3.2 Data charts

82 3.3 Methods

83 3.3.1 RQ1

84 3.3.2 RQ2

85 4 Evaluation

86 Evaluation.

87 5 Conclusions

88 Conclusions.

89 5.1 Acknowledgements

90 Many thanks.

91 References

- 92 [1] Alexis Conneau, Douwe Kiela, Holger Schwenk, Loïc Barrault, and Antoine
93 Bordes. Supervised learning of universal sentence representations from
94 natural language inference data. In *Proceedings of the 2017 Conference*
95 *on Empirical Methods in Natural Language Processing*, pages 670–680,
96 Copenhagen, Denmark, September 2017. Association for Computational
97 Linguistics.
- 98 [2] Zihang Dai, Zhilin Yang, Yiming Yang, Jaime G. Carbonell, Quoc V. Le, and
99 Ruslan Salakhutdinov. Transformer-xl: Attentive language models beyond
100 a fixed-length context. *CoRR*, abs/1901.02860, 2019.

- 101 [3] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert:
102 Pre-training of deep bidirectional transformers for language understanding.
103 *arXiv preprint arXiv:1810.04805*, 2018.
- 104 [4] Xiaodong Liu, Pengcheng He, Weizhu Chen, and Jianfeng Gao. Multi-task
105 deep neural networks for natural language understanding. *arXiv preprint*
106 *arXiv:1901.11504*, 2019.
- 107 [5] Matthew E. Peters, Mark Neumann, Mohit Iyyer, Matt Gardner,
108 Christopher Clark, Kenton Lee, and Luke Zettlemoyer. Deep contextualized
109 word representations. *CoRR*, abs/1802.05365, 2018.
- 110 [6] Alec Radford, Jeff Wu, Rewon Child, David Luan, Dario Amodei, and Ilya
111 Sutskever. Language models are unsupervised multitask learners. 2019.
- 112 [7] Google Trends. Explore. [https://trends.google.nl/trends/explore?](https://trends.google.nl/trends/explore?date=today%205-y&q=fake%20news)
113 [date=today%205-y&q=fake%20news](https://trends.google.nl/trends/explore?date=today%205-y&q=fake%20news), 2019. Retrieved on 16th of April, 2019.
- 114 [8] William Yang Wang. ”liar, liar pants on fire”: A new benchmark dataset
115 for fake news detection. *CoRR*, abs/1705.00648, 2017.